

Voicefork

Cloud Computing course project
A.A. 2022-2023

Faculty of Ingegneria dell'informazione, informatica e statistica
Department of Informatica



SAPIENZA
UNIVERSITÀ DI ROMA

Participants:


Corsi Danilo - 1742375

Lucciola Alessio - 1823638

Scarcelli Domiziano - 1872664

Introduction

- **Task:** Build a mobile application that allows users to manage their reservations at restaurants.
- **Goal:** Realize a structure based on several **microservices**, so that the application is **highly maintainable** and **scalable**.




Voicefork allows users to make online reservations at restaurants

Design of the solution

- Backend structured in different microservices
- Each microservice it's independent (has its own purpose and database)
- API built with **REST** technology used for single, well-defined tasks
- Additional technologies:
 - MinIO: To store any non relational data such as images
- **Docker**: Used in the local development environment



Minio Logo



Docker Logo

Design of the solution - Microservices

We implemented 5 microservices:

- **Users:** It handles users operations (registration and login)
- **Reservations:** It handles reservations operations (creation, visualization)
- **Restaurants:** It handles restaurants operations (searching)

Design of the solution - Microservices

- **Embeddings:** Used when searching for restaurants (encode restaurants characteristics, returns top k restaurants whose name is the most similar from the query)



[127, 123, 184, 543]

↑ ↓ L2 Distance

[124, 128, 184, 135]


Pizzeria da Triticum

Levenshtein Distance ↑ ↓

Pizza triticum

Average final distance

FAISS
Scalable Search With Facebook AI



Design of the solution - Microservices

- **Nginx:** It connects all the microservices under the same URL

`https://voicefork-api.com/`

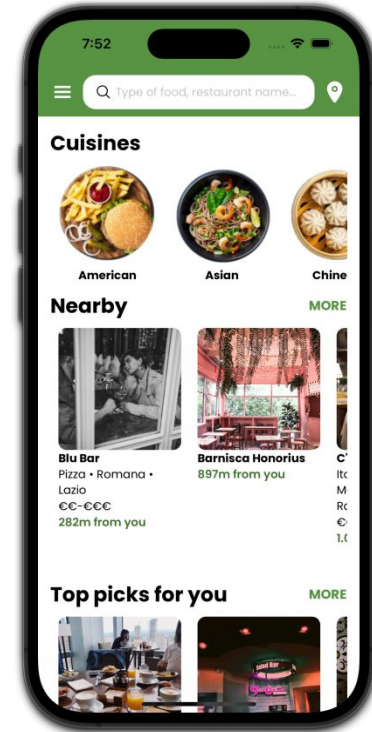
`https://voicefork-api.com/restaurants/`

`https://voicefork-api.com/users/`

`https://voicefork-api.com/reservations/`

Frontend Application

- Built with **React Native**
- Cross-platform application available on **Android** and **IOS**
- The main features are:
 - Login and registration
 - Searching for restaurants
 - Making reservations
 - Managing reservations
- Each microservice is involved in a certain task




Voicefork homepage

Frontend Application - Login and Registration


**Discover and
book the best
restaurants**

Discover and book the best restaurants
and leave reviews by also using your
voice

Email:

 Email

Password:

 Password

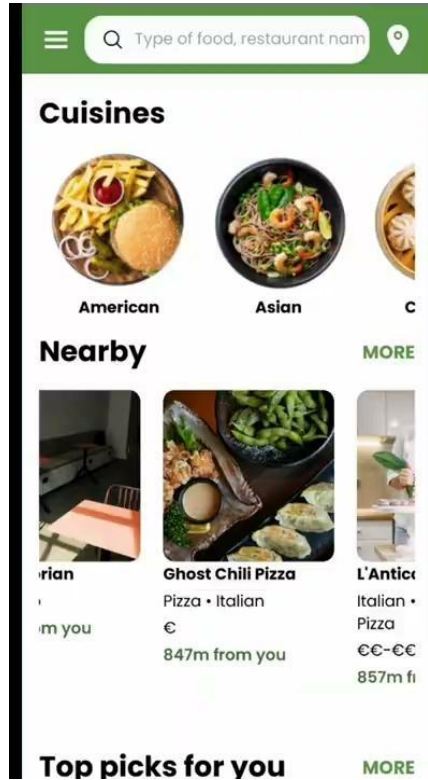
Login

Not a member yet? [Register Now](#)

Frontend Application - Searching and reservation



Frontend Application - Reservation list



Deployment

- Backend deployed on AWS:
 - **AWS ECS:** To deploy the microservices
 - **Task definition:** To replicate the docker configurations
 - Each microservice has its own container
 - **AWS Fargate:** To run containers
 - **AWS S3:** To store the unstructured data (images)
 - **MinIO:** To interface with AWS S3
 - **Amazon RDS:** To run the databases
- **Terraform:** To declare the entire cloud infrastructure through a set of files and initialize it with a simple command.



Amazon ECS



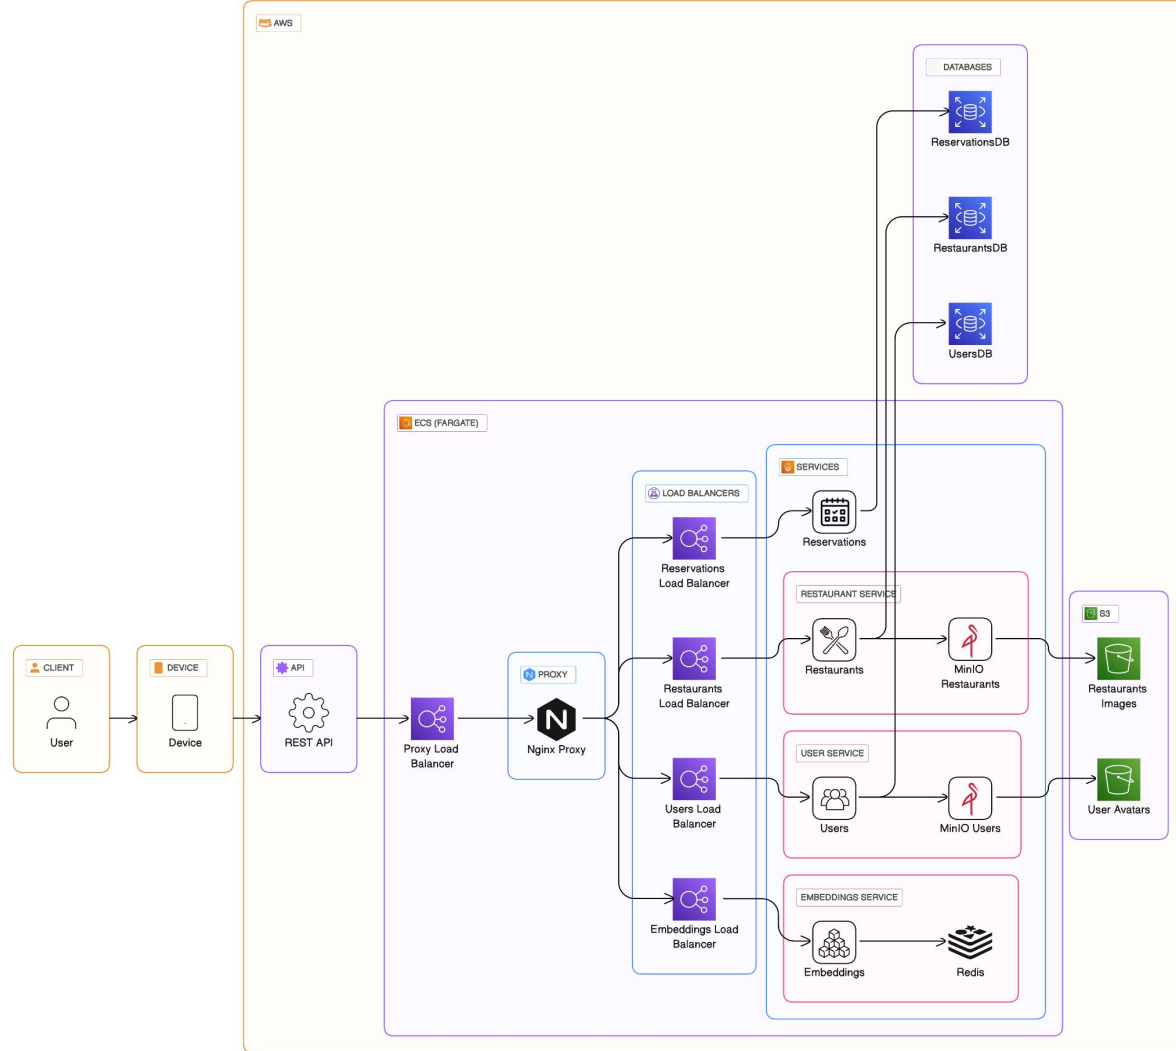
Amazon Fargate



Amazon S3




Amazon RDS



Real Scenario Experiments

- In order to simulate the application in a real scenario we used real restaurants:
 - “Tripadvisor European Restaurant” dataset
 - More than 200k Italian restaurants
 - Each restaurant contains real information such as name, address, average rating, type of cuisine, etc..
 - Dataset plugged into the database hosted on Amazon RDS




Tripadvisor Logo


Dataset source: <https://www.kaggle.com/datasets/stefanoleone992/tripadvisor-european-restaurants>

Design of experiments

- **K6:** tool that enables setup and run load tests
 - Simulate scenarios and application behavior
 - Provides real-time metrics and reports



- **Pipeline phases**



- **Autoscaling target**

- Min capacity: 1
- Max capacity: 10

- **Microservices specs**



512 units (0.5 vCPU)
(Embeddings - 1 vCPU)

+ more resources



1024MB
(Embeddings - 4096MB)

+ more resources




Load balancer
(distribute traffic)

- **Autoscaling policy**

- Target value: 60% (average percentage of CPU utilization)
- Scale-out cooldown: 120 seconds
- Scale-in cooldown: 180 seconds

Design of experiments - Scenario

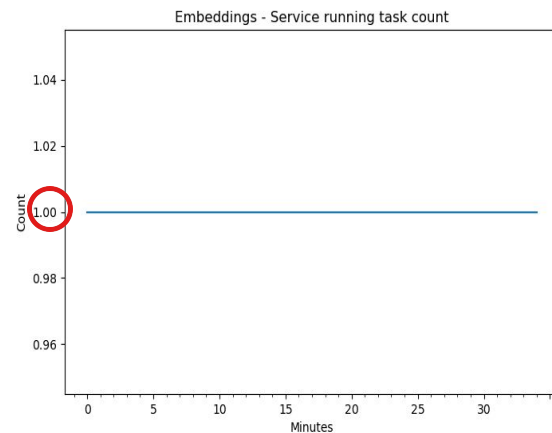

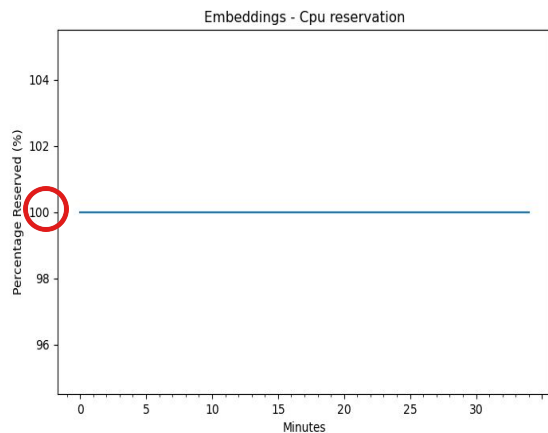
- The user has a real interaction with the application (simulate the behavior of the actual system)
- **Example:**



- All actions are randomized (users can perform any action at any time)
- Some actions are constrained by others (e.g. login before making reservation)
- A total of 10 tests were conducted (each lasting 30 minutes), the results were averaged

Results of experiments - Embeddings

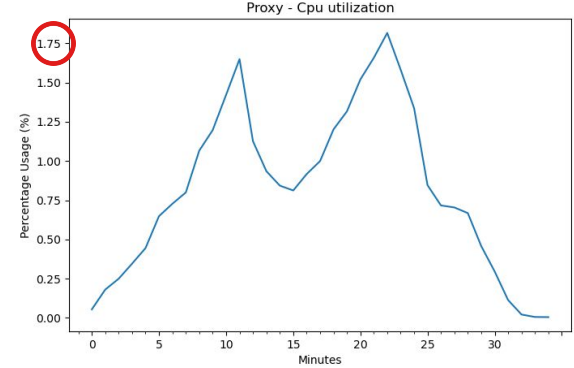
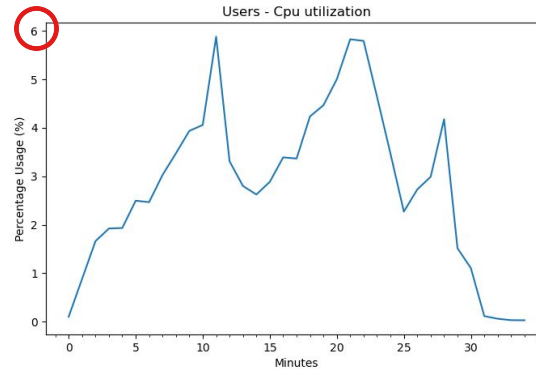
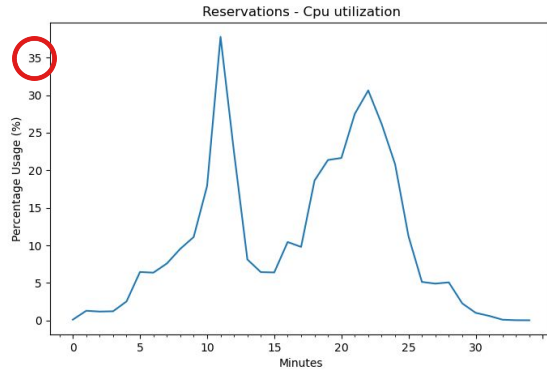
- No scaling operation (CPU utilization never exceeds 60%)
- CPU and memory reservations are constants (as well as the running task count)



- **“Reservation”** refers to the portion of memory and cpu that is allocated for each individual task

Results of experiments - Other microservices

- Have a similar trend (they don't perform any computationally heavy operations even if under a substantial load)




Results of experiments - Restaurants


Results of experiments - Restaurants

- There is a sudden increase with a peak of 80% of cpu utilization (target value was 60%)


Restaurants - Cpu utilization



Restaurants - Service running task count

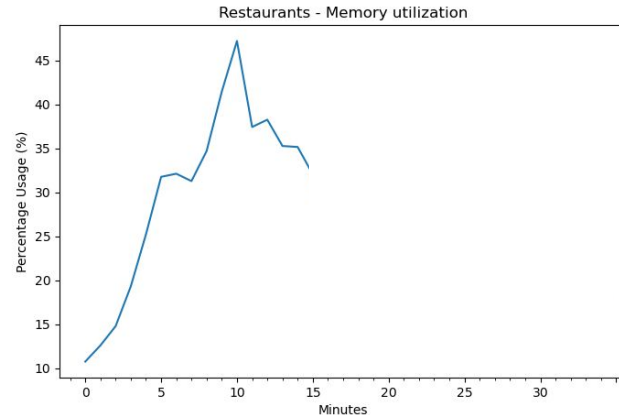

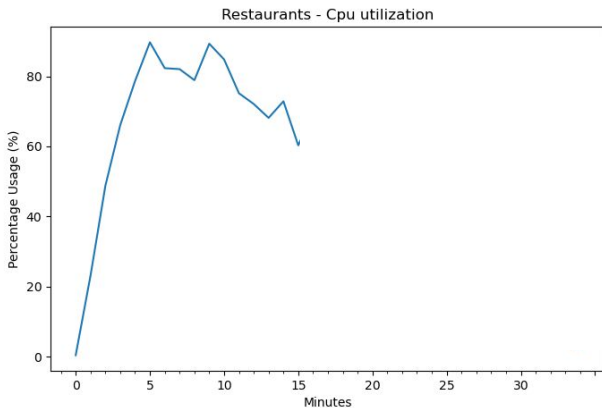


Restaurants - Memory utilization



Results of experiments - Restaurants


- There is a sudden increase with a peak of 80% of cpu utilization (target value was 60%)
- After the scale-out cooldown the scaling system starts working by adding new running tasks




Results of experiments - Restaurants

- There is a sudden increase with a peak of 80% of cpu utilization (target value was 60%)
- After the scale-out cooldown the scaling system starts working by adding new running tasks
- The usage percentage starts decreasing slowly


Restaurants - Cpu utilization



Restaurants - Service running task count

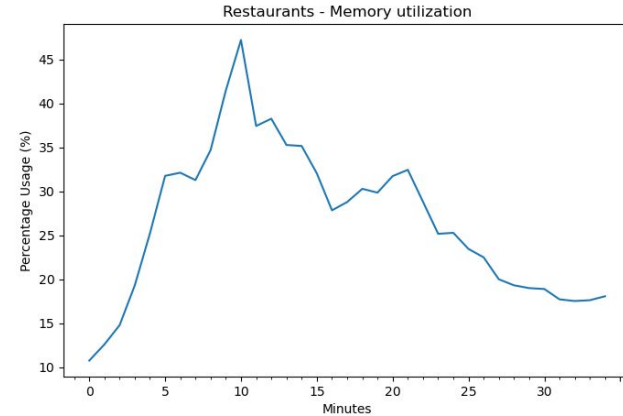

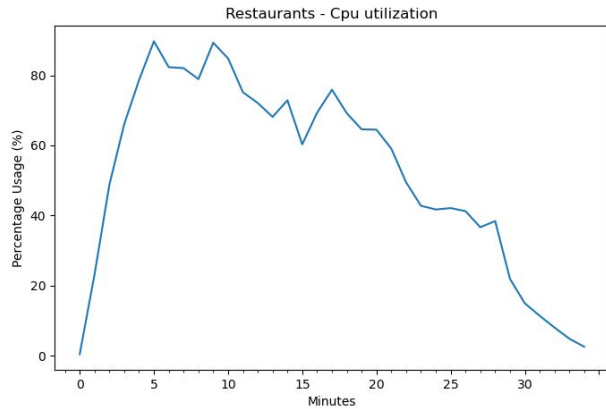


Restaurants - Memory utilization



Results of experiments - Restaurants

- There is a sudden increase with a peak of 80% of cpu utilization (target value was 60%)
- After the scale-out cooldown the scaling system starts working by adding new running tasks
- The usage percentage starts decreasing slowly
- After the rump-down period also the running tasks number decreases




Conclusions




Conclusions

- We built the application using the **microservices architecture**





Conclusions




- We built the application using the **microservices architecture**
- Taking advantages of **various benefits**
 - Independently deployable (we built and deployed each microservice in a parallel way)
 - Loosely coupled (each microservice is completely independent)
 - Organized by a small team (each member of the group focused on a specific service)

Conclusions

- We built the application using the **microservices architecture**
- Taking advantages of **various benefits**
 - Independently deployable (we built and deployed each microservice in a parallel way)
 - Loosely coupled (each microservice is completely independent)
 - Organized by a small team (each member of the group focused on a specific service)
- Each service was containerized using **Docker Compose**
 - Enabling fast deployment on **Amazon ECS**
 - Infrastructure managed by **Terraform**



Conclusions



- We built the application using the **microservices architecture**
- Taking advantages of **various benefits**
 - Independently deployable (we built and deployed each microservice in a parallel way)
 - Loosely coupled (each microservice is completely independent)
 - Organized by a small team (each member of the group focused on a specific service)
- Each service was containerized using **Docker Compose**
 - Enabling fast deployment on **Amazon ECS**
 - Infrastructure managed by **Terraform**
- Tests on the deployed application demonstrated the cloud's ability to manage workload efficiently (automatically adding and removing resources based on chosen policy)



Thank you for your attention